Economically-Efficient Data Stream Analysis

Roberto Oliveira Jr.
Advisor: Adriano Veloso
Co-advisor: Wagner Meira Jr.

Computer Science Dept - UFMG - Brazil

Data Stream

- Definition
  - Fast and possible unbounded sequence of data that arrives at time-varying.
- Motivation
  - It allows us to process huge volumes of data.
- Problem
  - Automatically extraction of relevant patterns and relations from data that is continuously created.
    - Keep track of data streams is useful for systems monitoring, online social network advertising, etc.

Social Networks Streams and Advertising

Superbowl 2013

Classification in Data Streams

- Classification models are applied to distinguish between pre-defined labels.

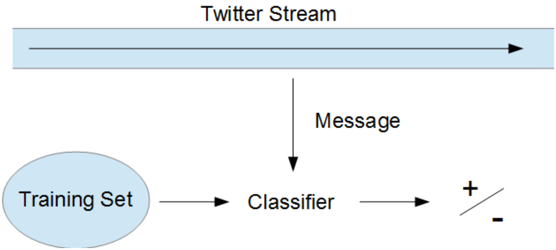

Classification in Data Streams

- Classification models are applied to distinguish between pre-defined labels.

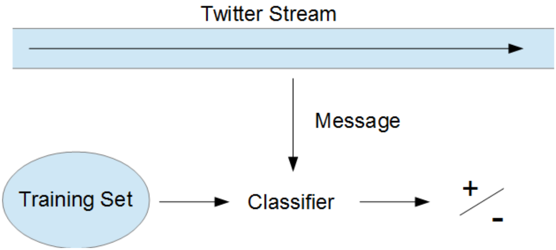

- Data characteristics may change with time.

Concept Drifts

- Concept Drift is unforeseen changes in data's nature over time.

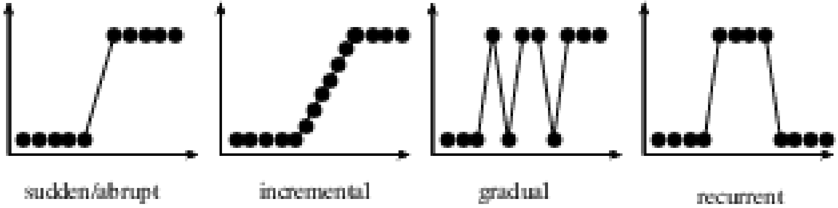

Concept Drifts

- Concept Drift is unforeseen changes in data's nature over time.

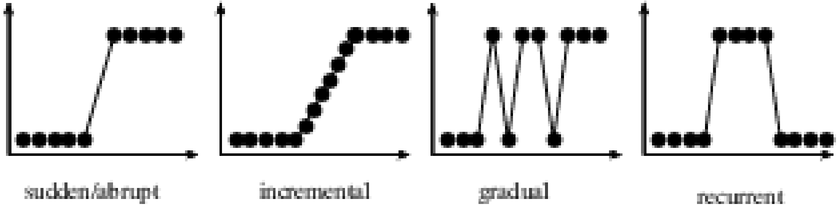

- Data streams contains combination of such patterns.

Sports (WC 2010)

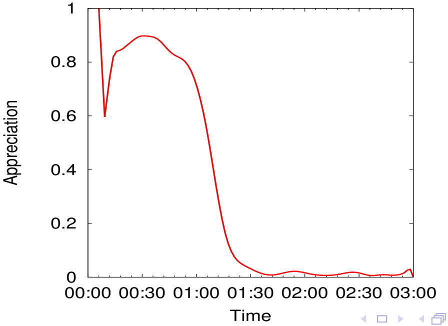

Elections (Brazil 2010)

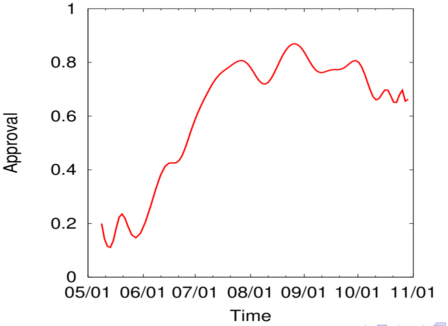

Classifying Data Streams

- Effective classification requires:
  - Updating the classification model as the stream evolves.
    - Taking into account resources limitation: memory, time and learning requirements.

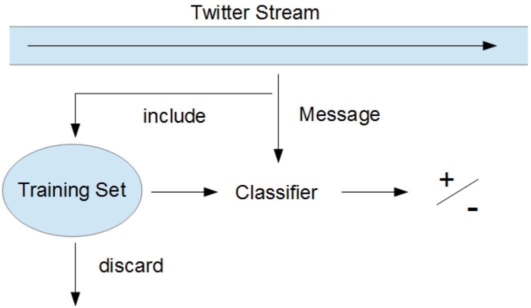

Research Question

How to deal with concept drifts?

1. Which classification model choose?
2. How to reduce labeling efforts?

Research Question

How to deal with concept drifts?

1. Which classification model choose?
2. How to reduce labeling efforts?

Classification Model

- Classification models are composed by association rules.
  - $\{x \to y\}$, where $x \in X$ and $y \in Y$
- Models are built on-the-fly:
  - For a given $[x_i, *]$, rules $\{x \to y\}$ such that $x \in x_i$ are produced.
  - Prediction is performed from the combination of these rules.
- At each time step is produced a model $\mathcal{R}(x_i)$.

Research Question

How to deal with concept drifts?

1. Which classification model choose?
2. How to reduce labeling efforts?

Dealing with Drifts

- Two properties are necessary in order to produce classifiers that are robust to drifts:
  - Adaptiveness:
    - The ability to adapt itself to drifts.
  - Memorability:
    - The ability to recover itself from drifts.

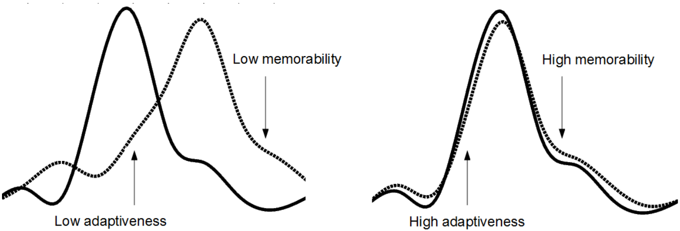

Dealing with Drifts

- Two properties are necessary in order to produce classifiers that are robust to drifts:
  - Adaptiveness:
    - The ability to adapt itself to drifts.
    - The training-set must contain fresh messages.
  - Memorability:
    - The ability to recover itself from drifts.
    - The training-set must contain pre-drift messages.
- Improving both properties simultaneously may lead to a conflict-objective problem.
  - Improve adaptiveness may hurt memorability, and vice-versa.

Pareto Efficiency

Example: hotels in Petrópolist.

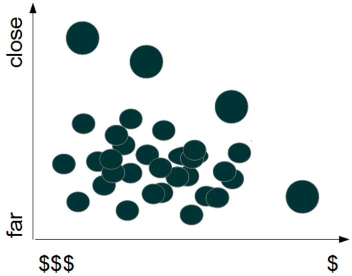

Pareto Efficiency

Pareto frontier: "when some action could be done to make someone better off without hurting anyone else, then it should be done."

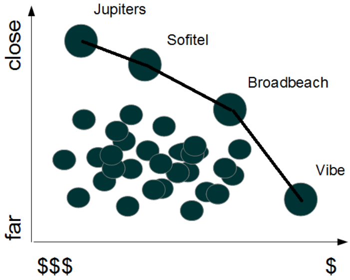

Compensation — Kaldor-Hicks Principle

Region of compensation: “when some action could be done to make someone better off, and this could compensate those that are made worse off, then it should be done.”

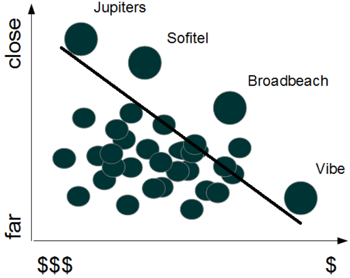

Utility Measures

- Distance in space:
  - How similar message t_j is to the newest message t_n.
  - $U_s(t_j) = \frac{|\mathcal{R}(t_n) \cap \mathcal{R}(t_j)|}{|\mathcal{R}(t_n)|}$
- Distance in time:
  - How fresh is the message.
  - $U_t(t_j) = \frac{\gamma(t_j)}{\gamma(t_n)}$.
    - $\gamma(t_j)$ returns the time in which message t_j arrived.
- Random permutation of messages:
  - $U_r(t_j) = \frac{\alpha(t_j)}{|\mathcal{D}_n|}$
    - $\alpha(t_j)$ returns the position of t_j in the shuffle.
    - $\mathcal{D}_n$ is the training set at time step n.

Utility Measures

1. At each time step *n*:
    1. Place candidate messages in the utility space.
    2. Select messages in the Pareto frontier.

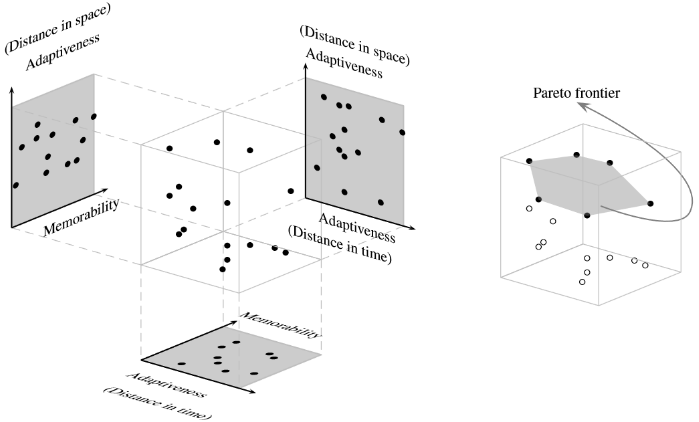

Research Question

How to deal with concept drifts?

1. Which classification model choose?
2. How to reduce labeling efforts?

Reducing Labeling Efforts

Evaluation

- Measures used:
  - Mean Squared Error.
  - RAM-Hours:
    - A GB of RAM deployed for 1 hour execution.
- Labeling Effort:
  - Different batch sizes and δ values.
  - Baselines:
    - AC – Active Classifiers (KDD 2011)
    - HAT – Hoeffding Adaptive Trees (JMLR 2011)
    - ILAC – Incremental Lazy Classifiers (SIGIR 2011)

Datasets

	Concept Drift Pattern			
Dataset	Sudden	Incremental	Gradual	Recurrent
Presidential Elections	-	X	X	-
Person of the Year	-	X	X	-
FIFA World Cup - EN	X	-	-	-
FIFA World Cup - PT	X	-	-	-
Cover Type	X	-	X	X
Spam Filtering	X	-	X	X
Poker Hand	-	-	X	X

Evaluation

- MSE and RAM-Hours

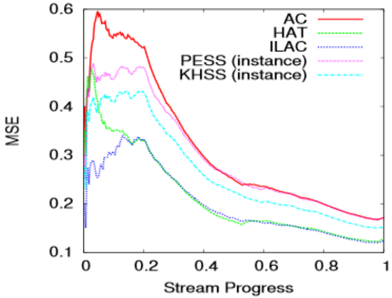

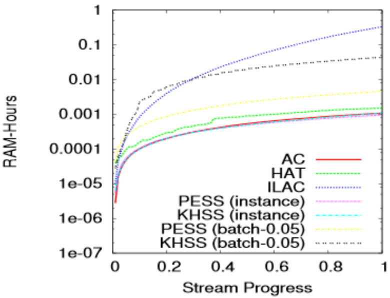

Evaluation

- MSE and Labeling Effort

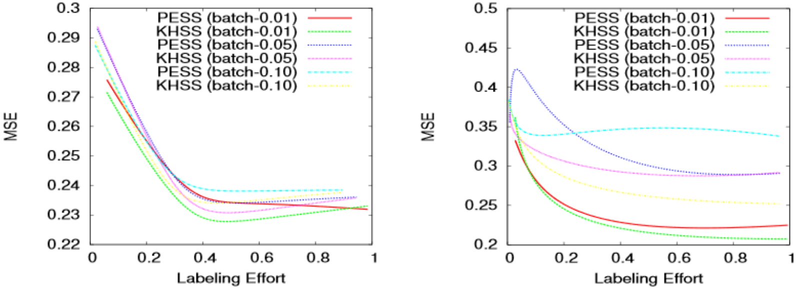

Conclusions

- Sentiment analysis on Twitter streams.
  - Limited computing and training resources.
  - Sentiment drifts.
- Efficiency and accuracy.
  - Incremental classifiers.
  - Pareto efficiency and compensation principle.
- Our results.
  - 50% reduction in terms of labeling effort without impact on accuracy.
- Future work includes:
  - Other utility measures.
  - Other application scenarios.

Thank you!

adrianov@dcc.ufmg.br